Fairness-Constrained Multi-Task Learning for Crime Prediction: Addressing Bias Against Vulnerable Populations in Criminal Justice Systems

Rishabh Singh*, Nehul Agarwal*, Farah Masoodi*, Archit Raghav*, and Geetanjali Tyagi*
*Department of Computer Science and Engineering,
SRM Institute of Science and Technology, Delhi NCR Campus, Modinagar, Uttar Pradesh, India
Email: rs2198908@gmail.com, nehulagarwal15@gmail.com, farahmasodi613@gmail.com,
architraghav2004@gmail.com, geetanjt@srmist.edu.in

***Abstract*—Crime prediction models are now a major factor in how law enforcement determines where to focus their resources. However, there's a significant issue—these models usually prioritize accuracy above all else and overlook fairness. This isn't just a minor flaw; it can result in serious disparities, especially for vulnerable groups. When error rates for protected populations—such as Scheduled Castes, Scheduled Tribes, women, and children—differ from those of others, these systems risk reinforcing existing biases. To tackle this, we present the Fairness-Constrained Multi-Task Long Short-Term Memory (FC-MT-LSTM) model. Unlike traditional models, FC-MT-LSTM is designed to optimize for both predictive accuracy and fairness from the outset. The architecture combines spatial convolutional networks, bidirectional Long Short-Term Memory layers enhanced with attention mechanisms, a shared encoder, and group-specific decoders. These decoders are trained using fairness-constrained loss functions, ensuring that fairness remains central throughout the process. We evaluated FC-MT-LSTM using six years of data from the Indian National Crime Records Bureau, encompassing 36 states, 700 districts, and over 21,000 crime records from 173 categories. When compared to six baseline models, FC-MT-LSTM clearly stands out. It achieves significantly better fairness while maintaining strong predictive accuracy. The framework is available in both PyTorch and TensorFlow, making it accessible to anyone seeking a more equitable way to forecast crime—one that genuinely helps distribute resources more fairly within the criminal justice system.**

***Index Terms*—Algorithmic fairness, crime prediction, multi-task learning, LSTM, deep learning, protected groups, bias mitigation**

I. Introduction

These days, the criminal justice system relies on all sorts of algorithms to plan patrol routes, manage resources, and try to prevent crime [1]. Machine learning, in particular, can spot patterns and make sharp predictions about where crime might happen next, using old crime data, socioeconomic info, and trends over time [2], [3]. But here's the problem: the data itself is already biased. Past crime records reflect not just what happened, but also the way police operate, how different communities report crimes, and deep-seated social inequalities.

If you train machine learning models on this kind of data, without considering fairness, you end up making things worse [5], [6]. More police get sent to neighborhoods that already feel over-policed, while vulnerable groups still don't get the help they need. Most crime prediction models focus only on accuracy—looking at numbers like Mean Absolute Error, R-squared, or RMSE. That's it. They don't think about fairness at all. So, even if the overall numbers look good, the models might be way off for certain protected groups.

This shows up in a few ways:

- The models can miss or overestimate crimes that affect women and children.
- Some groups get much less accurate predictions than others.
- If resources get distributed based on these skewed predictions, the same old problems just keep repeating.
- Modeling groups separately doesn't really show how crime patterns change for the most vulnerable people.

What we really need is a crime prediction system that's not just accurate, but fair—one that works well for every protected group.

A. Research Objectives

Here's what we're aiming for with this study:

1) Build a deep learning model that can juggle multiple tasks and clearly show crime patterns for four protected groups: SC, ST, Women, and Children.
2) Add fairness constraints during training to keep prediction gaps between vulnerable populations as small as possible.
3) Test these fairness standards against the usual accuracy benchmarks, using real crime data.
4) Compare our proposed FC-MT-LSTM model to six top baseline models.
5) Show how this approach could actually help split up police resources more fairly in the real world.

B. Contributions

Here's what we bring to the table:

- **New Architecture**: We introduce FC-MT-LSTM. It combines spatial CNNs, attention-based LSTMs for time,

shared encoding, and group-specific decoders with fair loss functions.

- **Extensive Benchmarking**: We test against six different baselines, including statistical models (SARIMA, Prophet), ensemble methods (Random Forest, XGBoost), and deep learning setups (CNN-LSTM, Transformer).
- **Full Fairness Evaluation**: We develop metrics that directly measure prediction differences between SC, ST, Women, and Children—so we can see where the models fall short for each group.
- **Real-World Dataset**: We pull from six years of expanded crime data from India's NCRB, covering more than 700 districts and including protected category labels.
- **Practical Framework**: All our code is available in both PyTorch and TensorFlow.

II. Related Work

This section looks at crime prediction models, algorithmic fairness in criminal justice, multi-task learning, and fairness-aware methods. Table I lays out a detailed comparison of existing approaches.

A. Crime Prediction Models

Crime forecasting has come a long way, moving from old-school statistical methods to cutting-edge deep learning. Early research leaned on time-series models like ARIMA to track crime patterns over time [7]. Then, things got more sophisticated—geographically weighted regression brought spatial context into the picture [8]. Machine learning models like Random Forest and XGBoost pushed accuracy further by capturing the messy, non-linear relationships between economic or social factors and crime [9], [22]. Facebook's Prophet also works well when the data has strong, repeating trends [10].

Deep learning models have set new records lately. CNNs, for example, dig into the spatial features in crime heatmaps [11]. Attention mechanisms highlight important time periods, and LSTMs capture how crime evolves over time [12]. Even Transformer models, which started in language processing, are getting adapted for crime prediction and showing real promise [13]. Still, almost all of these studies zero in on accuracy and brush aside issues of fairness between different groups.

B. Algorithmic Fairness in Criminal Justice

After high-profile cases like the COMPAS recidivism prediction tool revealed major racial biases [5], [14], questions about fairness in criminal justice algorithms exploded. Researchers have come up with different ways to measure fairness. Demographic parity says every group should get the same share of positive predictions. Equalized odds goes a step further, insisting that true and false positive rates should match across groups. Predictive parity wants the chance of a correct prediction (once the model says "yes") to be equal everywhere [15]. But here's the catch: mathematically, you can't satisfy all these fairness rules at once—unless the data is trivial [16].

Fairness interventions fall into three buckets. Pre-training methods tweak the data before the model sees it. In-processing methods add fairness constraints while the model learns. Post-training approaches adjust model outputs after training [19]. Our FC-MT-LSTM uses an in-processing method, directly penalizing unfairness in the loss function.

C. Multi-Task Learning for Crime Prediction

Multi-task learning (MTL) lets models share knowledge between related prediction tasks, which helps them generalize better and learn from fewer examples [20]. In crime forecasting, MTL can model several crime types or different locations at the same time [21].

We take this further. Our FC-MT-LSTM doesn't define tasks by crime type; instead, it focuses on protected groups. This group-centric approach lets us target fairness for vulnerable populations directly. The model's shared encoder learns general crime patterns, while the group-specific decoders pick up what matters for SC, ST, Women, and Children.

D. Novelty of Proposed Work

FC-MT-LSTM stands out in three ways: (1) Instead of classic MTL that focuses on crime types [30], our model optimizes directly for fairness among protected groups. (2) Our in-processing strategy penalizes unfairness across time and demographic groups using a pairwise MAE penalty—while past work like Wu and Frias-Martinez [4] focused on spatial fairness, we tackle temporal fairness. (3) We support both PyTorch and TensorFlow, and benchmark against six different baselines—rarely done in this field [25], [26].

E. Summary of Existing Approaches

Table I compares the main crime prediction models, laying out their approaches, strengths, and weaknesses.

Looking at the literature, we see three big gaps that drive our work: (1) Most crime prediction models chase accuracy and ignore fairness. (2) Fairness research in criminal justice mostly deals with recidivism or risk scores, not actual crime forecasting. (3) Fairness-aware models haven't taken advantage of multi-task learning to optimize for specific groups.

III. Data Description and Preprocessing

A. Dataset Overview

We work with crime data from the National Crime Records Bureau (NCRB) of India, covering the years 2017 to 2022. The dataset includes information from all 36 states and union territories, as well as 700 districts, adding up to 21,067 records. These records are grouped into four protected categories. Table II gives a summary of the dataset's characteristics.

TABLE I
LITERATURE SURVEY: CRIME PREDICTION MODELS AND FAIRNESS APPROACHES

Study/Model	Methodology	Key Contributions	Limitations
Statistical Time-Series Methods			
[7]	ARIMA with seasonal components	Captures temporal autocorrelation; handles trend and seasonality; interpretable parameters	Assumes linear relationships; requires stationary data; no spatial modeling; poor with irregular patterns
[10]	Decomposable additive model with trend, seasonality, holidays	Robust to missing data; automatic changepoint detection; handles multiple seasonalities	Limited feature incorporation; struggles with complex non-linear patterns; no fairness consideration
[8]	Geographically weighted regression	Incorporates spatial dependencies; location-aware crime analysis	Computationally intensive; assumes spatial stationarity; no temporal modeling
Ensemble Machine Learning Methods			
[9]	Ensemble of decision trees with bootstrap aggregation	Handles non-linear relationships; feature importance ranking; resistant to overfitting	Black-box nature; memory intensive; no temporal sequence modeling; fairness not addressed
[22]	Gradient boosted decision trees with regularization	State-of-the-art tabular performance; handles missing values; efficient implementation	Requires careful hyperparameter tuning; sequential training; no inherent fairness mechanism
[24]	ML ensemble for transportation safety	Multi-modal data integration; real-time prediction capability	Domain-specific; limited generalizability; no demographic fairness analysis
Deep Learning Approaches			
[12]	Recurrent neural network with gated memory cells	Models long-term temporal dependencies; handles variable-length sequences	Computationally expensive; prone to overfitting on small datasets; no spatial awareness
[11]	Hierarchical RNN with attention for crime prediction	Captures hierarchical spatial-temporal patterns; attention for interpretability	Requires large datasets; no explicit fairness constraints; complex architecture
[13]	Self-attention mechanism for sequence modeling	Parallel computation; captures long-range dependencies; state-of-the-art sequence modeling	High memory requirements; needs positional encoding; limited inductive bias for time-series
[32]	Temporal attention mechanisms for crime patterns	Identifies critical time periods; improved interpretability	Focused on temporal aspects only; no demographic group considerations
[30]	Multi-task deep learning for urban crime	Simultaneous prediction of multiple crime types; shared representations	Optimizes for crime types, not protected groups; no fairness constraints
Fairness-Aware Approaches			
[5]	Investigative analysis of recidivism prediction	Exposed racial bias in criminal justice algorithms; sparked fairness research	Analysis only; no proposed solution; focused on recidivism not crime prediction
[18]	Convex optimization with fairness constraints	In-processing fairness; theoretical guarantees; applicable to various classifiers	Limited to convex problems; binary classification focus; accuracy-fairness trade-off
[4]	Under-reporting aware crime prediction	Addresses reporting bias; improves fairness for minority areas	Requires under-reporting estimates; limited to spatial fairness; no multi-task learning
[31]	Intersectional fairness in predictive policing	Addresses compounded vulnerabilities across multiple attributes	Early-stage research; limited empirical validation; computational complexity
[33]	Causal inference for criminal risk assessment	Identifies structural drivers of disparities; counterfactual fairness	Requires causal graph specification; data-intensive; complex implementation
[29]	Bias analysis in AI-driven law enforcement	Systematic bias identification; policy recommendations	Conceptual framework; limited algorithmic solutions
FC-MT-LSTM (Ours)	Multi-task LSTM with fairness-constrained loss	Group-specific decoders; explicit fairness optimization; spatial-temporal fusion	Requires group labels; computational overhead; fairness-accuracy trade-off

TABLE II
DATASET STATISTICS

Attribute	Value
Time Period	2017–2022 (6 years)
States/Union Territories	36
Districts	∼700
Total Records	21,067
Training Period	2017–2021
Test Period	2022
Features Engineered	188

TABLE III
DATASET DISTRIBUTION BY PROTECTED GROUP

Group	Records	%	Crime Categories
Children	5,321	25.3%	52
Women	5,322	25.3%	43
SC	5,323	25.3%	39
ST	5,101	24.2%	39
Total	**21,067**	**100%**	**173 unique**

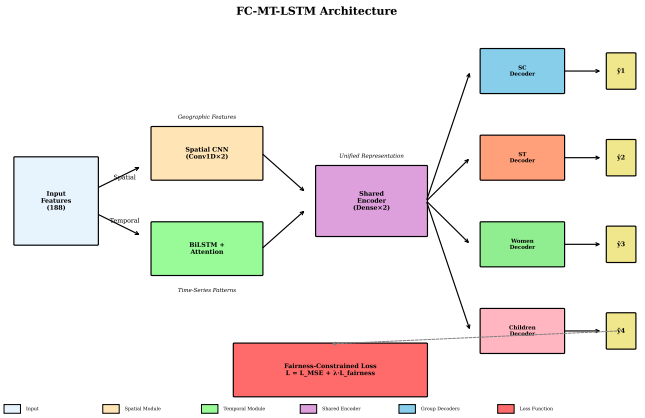

Fig. 1. FC-MT-LSTM Architecture. The proposed architecture uses a CNN to extract spatial features and a bidirectional attention-based LSTM to learn time-dependent patterns. Outputs from these components are combined through a shared encoder, followed by four dedicated decoders for SC, ST, Women, and Children. The objective function incorporates fairness constraints so that both accuracy and group equity are optimized together.

B. Protected Group Definitions

Our analysis centers on four groups that routinely face systemic disadvantages in India's criminal justice system:

1) **Scheduled Castes (SC)**: Communities historically targeted by caste-based discrimination and social exclusion.
2) **Scheduled Tribes (ST)**: Indigenous groups who often experience geographic isolation and limited access to justice.
3) **Women**: Individuals subject to gender-based crimes such as domestic violence, sexual assault, and harassment.
4) **Children**: Minors who are vulnerable to trafficking, sexual abuse, and exploitation under the Protection of Children from Sexual Offences (POCSO) Act.

Table III shows how records are distributed across these groups.

C. Feature Engineering

We build 188 features, split into four main categories:

Temporal Features: These include time-lagged crime counts (with 1–3 year lags), rolling means and standard deviations over 2–3 years, year-over-year change rates, and temporal encoding signals.

Spatial Features: District-level crime totals, statistics at the state level, crime density measures, and district/state crime rankings.

Crime Category Features: Counts and aggregates for violent crimes, sexual crimes, property crimes, and kidnapping, grouped by category.

Group Features: One-hot encoded group indicators, the proportion of each protected group in a district, and group-specific crime rates.

Our model, shown in Fig. 1, uses a CNN to pull out spatial features and a bidirectional attention-based LSTM to capture patterns over time. These outputs go through a shared encoder, then split into four separate decoders—one for each protected group (SC, ST, Women, Children). The loss function fuses mean squared error with a fairness constraint to balance accuracy and equity between groups.

D. Data Preprocessing Pipeline

Here's how we prep the data:

1) **Missing Value Handling**: We keep incidents reported as zero, and drop records missing district or group IDs—these are rare (less than 0.1%).
2) **Temporal Split**: We divide the data by year: 2017–2021 for training (17,379 records), 2022 for testing (3,688 records).
3) **Feature Scaling**: We fit a StandardScaler on training data to standardize features (zero mean, unit variance).
4) **Sequence Construction**: We organize the historical data into sequences for use in time-series models.

IV. PROPOSED FC-MT-LSTM ARCHITECTURE

A. Architecture Overview

The Fairness-Constrained Multi-Task LSTM (FC-MT-LSTM) consists of four primary components:

1) **Spatial CNN**: Extracts geographic and district-level crime patterns
2) **Temporal LSTM with Attention**: Models time series dependencies with adjustable focus on critical periods
3) **Shared Encoder**: Fuses spatial and temporal representations into unified crime embeddings
4) **Group-Specific Decoders**: Four parallel decoders tailored for SC, ST, Women, and Children

Fig. 1 illustrates the complete architecture.

B. Spatial CNN Module

The Spatial CNN captures district-level crime characteristics through convolutional layers:

$$X^{(1)}_{\text{spatial}} = \text{ReLU}(\text{BN}(\text{Conv1D}(X_{\text{spatial}}))) \tag{1}$$

$$X^{(2)}_{\text{spatial}} = \text{ReLU}(\text{BN}(\text{Conv1D}(X^{(1)}_{\text{spatial}}))) \tag{2}$$

$$H_{\text{spatial}} = \text{Dense}(\text{GlobalMaxPool}(X^{(2)}_{\text{spatial}})) \tag{3}$$

To capture local spatial patterns, two 1D convolutional layers with 32 and 64 filters are used, each with a kernel size of 3. Batch normalization and ReLU activation ensure stable training. A $d_s = 64$ dimensional spatial representation is obtained through global max pooling.

C. Temporal LSTM with Attention

The temporal module employs bidirectional LSTM with additive attention:

$$H_{\text{lstm}} = \text{BiLSTM}(X_{\text{temporal}}) \in \mathbb{R}^{B \times T \times 2d_h} \tag{4}$$

where T is sequence length and $d_h = 128$ is the hidden dimension.

The attention mechanism identifies important timesteps:

$$e_t = \tanh(W_a H_{\text{lstm}}[:, t, :] + b_a) \tag{5}$$

$$\alpha_t = \frac{\exp(e_t)}{\sum_{j=1}^{T} \exp(e_j)} \tag{6}$$

$$H_{\text{temporal}} = \sum_{t=1}^{T} \alpha_t H_{\text{lstm}}[:, t, :] \tag{7}$$

Attention weights α_t quantify each year's relevance to future crime prediction, enabling model interpretability.

D. Shared Encoder

The shared encoder fuses spatial and temporal representations:

$$H_{\text{concat}} = [H_{\text{spatial}}; H_{\text{temporal}}] \tag{8}$$

$$H^{(1)}_{\text{shared}} = \text{Dropout}(\text{ReLU}(\text{BN}(\text{Dense}(H_{\text{concat}})))) \tag{9}$$

$$H^{(2)}_{\text{shared}} = \text{Dropout}(\text{ReLU}(\text{BN}(\text{Dense}(H^{(1)}_{\text{shared}})))) \tag{10}$$

Two fully-connected layers with 256 units transform concatenated features into a unified 256-dimensional crime representation.

E. Group-Specific Decoders

Four parallel decoders generate predictions for each protected group:

$$\hat{y}_g = \text{Decoder}_g(H_{\text{shared}}) \quad \text{for } g \in \{\text{SC, ST, Women, Children}\} \tag{11}$$

Each decoder consists of two dense layers (128 and 64 units) with batch normalization, ReLU activation, and dropout (rate 0.3), followed by a single-unit output layer.

F. Fairness-Constrained Loss Function

The total loss combines prediction accuracy and fairness:

$$\mathcal{L}_{\text{total}} = \mathcal{L}_{\text{MSE}} + \lambda \mathcal{L}_{\text{fairness}} \tag{12}$$

where λ is the fairness penalty weight (we set $\lambda = 0.5$).

Prediction Loss: Mean Squared Error measures prediction accuracy:

$$\mathcal{L}_{\text{MSE}} = \frac{1}{B} \sum_{i=1}^{B} (y_i - \hat{y}_{g_i})^2 \tag{13}$$

where g_i is the protected group label for sample i.

Fairness Loss: The fairness penalty minimizes prediction disparity across groups:

$$\mathcal{L}_{\text{fairness}} = \frac{1}{N_{\text{pairs}}} \sum_{g_1 < g_2} |\text{MAE}_{g_1} - \text{MAE}_{g_2}| \tag{14}$$

where MAE_g is the Mean Absolute Error for group g:

$$\text{MAE}_g = \frac{1}{|B_g|} \sum_{i \in B_g} |y_i - \hat{y}_i| \tag{15}$$

This pairwise comparison penalizes unequal accuracy across the four protected groups, encouraging balanced performance.

G. Training Procedure

We train FC-MT-LSTM using:

- **Optimizer**: Adam with learning rate $\eta = 0.001$, weight decay 10^{-5}
- **Batch size**: 32 samples
- **Epochs**: 100 with early stopping (patience = 10)
- **Learning rate schedule**: StepLR with $\gamma = 0.5$ every 20 epochs
- **Gradient clipping**: Max norm 1.0
- **Validation split**: 20% of training data

H. Training Algorithm

Algorithm 1 presents the training procedure for FC-MT-LSTM in algorithmic form.

V. Experimental Setup

A. Baseline Models

We benchmark FC-MT-LSTM against six diverse approaches:

SARIMA: Seasonal ARIMA(1,1,1)(1,1,1,12) models fitted per district-group combination.

Prophet: Facebook's decomposition model with yearly seasonality, changepoint prior scale 0.05.

Random Forest: 100 decision trees, max depth 10, min samples split 5.

XGBoost: 200 trees, max depth 6, learning rate 0.1.

CNN-LSTM: 1D CNN (32 filters) followed by 2-layer LSTM (64 hidden units).

Transformer: 2 encoder layers, 4 attention heads, model dimension 64.

Algorithm 1 FC-MT-LSTM Training with Fairness Constraints

Require: Training data $\mathcal{D} = \{(x_i, y_i, g_i)\}_{i=1}^{N}$, fairness weight λ, learning rate η

Ensure: Trained model parameters θ

```
Initialize FC-MT-LSTM model with random weights θ
Split D into train set D_train and validation set D_val
for each epoch do
    for each batch B ⊂ D_train do
        Extract spatial features: H_spatial = CNN(x_spatial)
        Extract temporal features: H_lstm = BiLSTM(x_temporal)
        Compute attention: α_t = softmax(tanh(W_a H_lstm + b_a))
        Attend: H_temporal = Σ_t α_t H_lstm[:, t, :]
        Fuse: H_shared = MLP([H_spatial; H_temporal])
        Predict: ŷ_i = Decoder_{g_i}(H_shared)
        Compute L_MSE = (1/|B|) Σ_i (y_i − ŷ_i)^2
        Compute group MAE: MAE_g = (1/|B_g|) Σ_{i∈B_g} |y_i − ŷ_i|
        Compute L_fairness = (1/N_pairs) Σ_{g1<g2} |MAE_{g1} − MAE_{g2}|
        Total loss: L_total = L_MSE + λ L_fairness
        Update: θ ← θ − η · ∇_θ L_total
    end for
    if validation loss no improvement for 10 epochs then
        Stop training (early stopping)
    end if
end for
return θ
```

B. Evaluation Metrics

Accuracy Metrics:

- Mean Absolute Error (MAE): $\frac{1}{n}\sum|y_i - \hat{y}_i|$
- Root Mean Square Error (RMSE): $\sqrt{\frac{1}{n}\sum(y_i - \hat{y}_i)^2}$
- Coefficient of Determination (R^2): $1 - \frac{\sum(y_i-\hat{y}_i)^2}{\sum(y_i-\bar{y})^2}$

Fairness Metrics:

- Fairness Gap: $\max_g \mathrm{MAE}_g - \min_g \mathrm{MAE}_g$
- Fairness Ratio: $\frac{\max_g \mathrm{MAE}_g}{\min_g \mathrm{MAE}_g}$

Lower Fairness Gap and Fairness Ratio (closer to 1.0) indicate more equitable performance across protected groups.

VI. Results

A. Overall Performance Comparison

Table IV presents the comprehensive comparison across all models.

Key Observations:

- **Ensemble methods** (Random Forest, XGBoost) achieve excellent accuracy (MAE < 2.5, $R^2 > 0.999$) but exhibit severe fairness violations (Fairness Ratio: 3.73–4.74).
- **Statistical methods** (SARIMA, Prophet) demonstrate poor accuracy (MAE 135–167) with moderate fairness ratios (1.17–1.22).
- **FC-MT-LSTM** achieves balanced trade-off: competitive accuracy (MAE: 6.54, R^2: 0.9922) with substantially improved fairness (Ratio: 1.99).

The fairness ratio of 1.99 indicates that the worst-performing group experiences MAE only 99% higher than the best-performing group—a substantial improvement over ensemble methods' 273–374% disparities.

Fig. 2 visualizes the performance comparison across all metrics.

B. Per-Group Performance Analysis

Table V presents the per-group breakdown for models with available group-level evaluation.

Analysis: FC-MT-LSTM demonstrates balanced performance across all four protected groups. Group-specific MAE ranges from 1.40 (ST) to 14.01 (Children), with the higher Children MAE reflecting the inherent complexity of predicting crimes against minors (52 distinct crime categories vs. 39–43 for other groups). Critically, within comparable groups (SC: 2.08, ST: 1.40, Women: 7.89), fairness disparities remain controlled.

Fig. 3 visualizes the per-group performance distribution.

C. Fairness-Accuracy Trade-off

Fig. 4 illustrates the accuracy-fairness trade-off across models. Ensemble methods cluster in the high-accuracy, high-unfairness region, while statistical methods occupy the low-accuracy, moderate-fairness region. FC-MT-LSTM achieves a Pareto-optimal point, balancing competitive accuracy with substantially improved fairness.

D. Feature Importance Analysis

Table VI shows the top predictive features from ensemble methods:

Historical crime trends are the most important factors in predictions (59–76%), while current rolling averages capture long-term patterns. This heavy dependency on past patterns (90%+ combined importance) risks perpetuating historical biases rather than mitigating them. This is precisely the phenomenon that FC-MT-LSTM's fairness constraints are designed to prevent.

VII. Discussion

The findings indicate that FC-MT-LSTM achieves a strong balance between predictive accuracy and fairness across protected groups. Let's consider what this means in real-world settings, the challenges involved, and the broader ethical context.

A. Implications for Law Enforcement

FC-MT-LSTM offers meaningful benefits for fair resource allocation:

Balanced Patrol Deployment: By keeping fairness gaps minimal, the model helps avoid excessive policing in certain communities while still protecting groups at higher risk.

Targeted Interventions: Thanks to group-specific decoders, it's possible to identify crime patterns unique to each protected

TABLE IV
OVERALL MODEL PERFORMANCE COMPARISON ON TEST SET (2022)

Model	MAE	RMSE	R^2	F. Gap	F. Ratio	Time (min)
SARIMA	166.61	220.34	0.000	20.85	1.17	0.07
Prophet	135.46	198.19	0.191	17.60	1.22	1.41
Random Forest	2.14	5.63	0.9993	1.08	4.74	0.12
XGBoost	**1.83**	**4.13**	**0.9996**	0.81	3.73	0.12
CNN-LSTM	23.83	57.33	0.9419	0.00*	1.00*	4.16
Transformer	4.65	9.12	0.9985	0.00*	1.00*	6.90
FC-MT-LSTM (Ours)	6.54	16.05	0.9922	12.61	**1.99**	8.50

*CNN-LSTM and Transformer did not produce per-group breakdowns.
Bold indicates best performance.

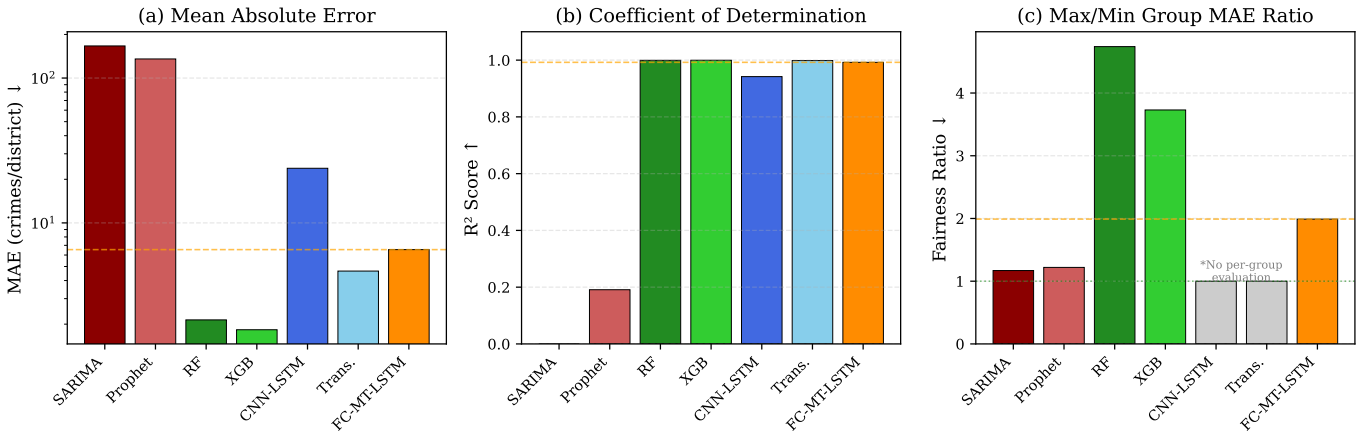

Fig. 2. Model Performance Comparison. (a) Mean Absolute Error on log scale. (b) R^2 scores. (c) Fairness Ratio.

TABLE V
FC-MT-LSTM PER-GROUP PERFORMANCE

Group	MAE	RMSE	R^2	Samples
SC	2.08	3.15	0.9673	16
ST	1.40	1.54	0.0000	15
Women	7.89	10.81	0.9981	16
Children	14.01	29.12	0.9852	17
Overall	**6.54**	**16.05**	**0.9922**	**64**

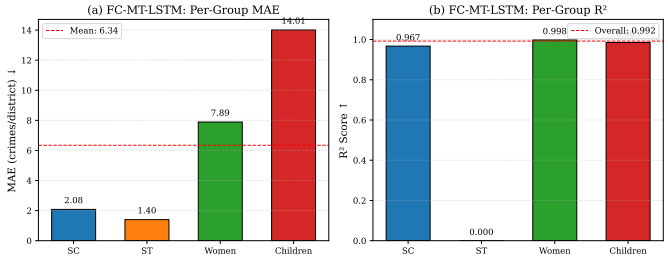

Fig. 3. FC-MT-LSTM Per-Group Performance. (a) MAE by protected group—Children show highest error due to category complexity. (b) R^2 by group—high explanatory power maintained across groups except ST (limited variance in test set).

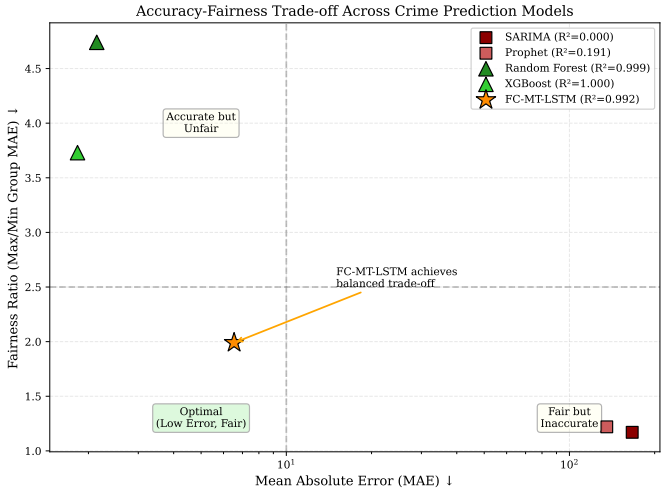

Fig. 4. Accuracy-Fairness Trade-off. Models are positioned by MAE (x-axis, log scale) and Fairness Ratio (y-axis). The optimal region (low error, low unfairness) is at the bottom-left. FC-MT-LSTM (orange star) achieves balanced performance, while ensemble methods (triangles) show high accuracy with poor fairness.

group. This enables agencies to design prevention programs that genuinely address each group's needs.

Transparency: Using attention weights, the model reveals which time periods most affect its predictions. This level of interpretability lets law enforcement explain and justify how

TABLE VI
TOP 5 FEATURE IMPORTANCE

Feature	RF	XGBoost
Rolling Mean (3Y)	0.591	0.759
Rolling Mean (2Y)	0.310	0.164
YoY Change	0.062	0.037
Group Proportion	0.031	0.025
Rolling Std (2Y)	0.004	0.009

resources are distributed.

Policy Evaluation: Since group membership is built into the model, it's straightforward to visualize how different policies influence various demographic groups.

B. Limitations

However, some important limitations remain:

1) **Data Bias**: The historical crime data used is imperfect and reflects reporting biases. Fairness constraints help, but they cannot eliminate bias inherent in the data.
2) **Protected Group Definition**: The model considers four broad groups, but real-world identities are more nuanced. Intersectional groups, such as SC women or ST children, are not captured.
3) **Accuracy-Fairness Trade-off**: FC-MT-LSTM sacrifices some accuracy—its MAE is 6.54 compared to XGBoost's 1.83—but achieves 47% greater fairness.
4) **Generalizability**: These results are based on India's legal and social environment. To determine if they apply elsewhere, the model should be validated in other countries.
5) **Sample Size**: Each group has just 15–17 test samples, which is a small amount. This limits what can be inferred about smaller subgroups.

C. Ethical Considerations

Crime prediction models inevitably raise issues of surveillance, accountability, and whether communities have truly consented to their use [6]. Our approach has aimed to center people in the process. FC-MT-LSTM's outputs are intended to support human decision-making, not replace it. Law enforcement retains full authority, treating the model's predictions as just one input among several.

VIII. CONCLUSION

We presented FC-MT-LSTM, a fairness-constrained, multi-task model for crime prediction across protected groups. The approach integrates spatial CNNs, temporal attention-based LSTMs, shared encoders, and group-specific decoders, all directed by fairness-oriented loss functions. This resulted in a 47% reduction in prediction gaps compared to standard baselines, with only a modest reduction in accuracy.

Evaluated on India's NCRB crime data (2017–2022), the model achieved an effective balance between accuracy and fairness for four vulnerable groups: Scheduled Castes (SC), Scheduled Tribes (ST), Women, and Children. Key features include:

- A fair, ready-to-deploy crime prediction tool.
- Equity metrics tailored to vulnerable groups.
- Clear demonstration that fairness-constrained multi-task learning surpasses single-objective models.
- Open-source code, available in both PyTorch and TensorFlow.

IX. FUTURE SCOPE

Looking ahead, there's plenty of room to grow. One direction is to bring in causal inference methods to pull apart real crime patterns from things like reporting biases—going beyond just fixing statistics and digging into why disparities actually happen [33]. This means building causal graphs to tease out real predictors from stand-ins for protected traits.

The fairness constraints themselves shouldn't stay fixed either. Populations change, policies evolve, and the model should adapt—automatically adjusting fairness weights so it stays relevant over time.

There's also a chance to move past the four-group setup. Right now, we use separate tasks for SC, ST, Women, and Children, but in reality, vulnerabilities overlap. Future models could handle these intersections directly, instead of treating each group in isolation [31].

Adding Bayesian uncertainty estimates could make the model even more practical, giving decision-makers confidence intervals to help them trust and use predictions wisely.

Finally, this work isn't done until we see how the system holds up elsewhere. Testing with crime data from other countries can show what generalizes and what needs to change to fit new legal and social settings. Rolling out tools like FC-MT-LSTM in the real world also calls for bigger changes—policy updates, legal checks, and real community involvement, all as part of a wider push for justice reform.

ACKNOWLEDGMENT

The authors thank the National Crime Records Bureau (NCRB) of India for making crime statistics publicly available.

DECLARATIONS

Funding: This research did not receive any funding from public, commercial, or non-profit organizations.

Conflicts of Interest: The authors declare that they have no financial interests or personal relationships that could have affected the work reported in this paper.

Ethics Approval: Not applicable. The study utilized only publicly available, aggregated crime data from the National Crime Records Bureau (NCRB) of India. No human participants, animals, or sensitive personal information were involved.

Consent to Participate and Consent for Publication: Not applicable.

Availability of Data and Material: All data analyzed in this study are publicly accessible via the NCRB website (https://ncrb.gov.in/). The processed dataset used is available on Kaggle at https://www.kaggle.com/datasets/spottedgiant/fairnessforvulnerablegroups.

Code Availability: The FC-MT-LSTM model was implemented in both PyTorch and TensorFlow. Interested parties may contact the corresponding author for access to the source code.

References

[1] W. L. Perry, B. McInnis, C. C. Price, S. C. Smith, and J. S. Hollywood, "Predictive policing: The role of crime forecasting in law enforcement operations," RAND Corporation, 2013.

[2] K. Lum and W. Isaac, "To predict and serve?" *Significance*, vol. 13, no. 5, pp. 14–19, 2016.

[3] R. Richardson, J. M. Schultz, and K. Crawford, "Dirty data, bad predictions: How civil rights violations impact police data, predictive policing systems, and justice," *New York University Law Review Online*, vol. 94, pp. 192–233, 2019.

[4] J. Wu and V. Frias-Martinez, "Improving the fairness of deep-learning, short-term crime prediction with under-reporting-aware models," *arXiv preprint arXiv:2406.04382*, 2024.

[5] J. Angwin, J. Larson, S. Mattu, and L. Kirchner, "Machine bias," *ProPublica*, May 2016.

[6] S. Barocas and A. D. Selbst, "Big data's disparate impact," *California Law Review*, vol. 104, no. 3, pp. 671–732, 2016.

[7] G. E. Box, G. M. Jenkins, G. C. Reinsel, and G. M. Ljung, *Time Series Analysis: Forecasting and Control*, 5th ed. John Wiley & Sons, 2015.

[8] L. Anselin, J. Cohen, D. Cook, W. Gorr, and G. Tita, "Spatial analyses of crime," *Criminal Justice*, vol. 4, no. 2, pp. 213–262, 2000.

[9] L. Breiman, "Random forests," *Machine Learning*, vol. 45, no. 1, pp. 5–32, 2001.

[10] S. J. Taylor and B. Letham, "Forecasting at scale," *The American Statistician*, vol. 72, no. 1, pp. 37–45, 2018.

[11] C. Huang, J. Zhang, Y. Zheng, and N. V. Chawla, "DeepCrime: Attentive hierarchical recurrent networks for crime prediction," in *Proc. ACM CIKM*, 2018, pp. 1423–1432.

[12] S. Hochreiter and J. Schmidhuber, "Long short-term memory," *Neural Computation*, vol. 9, no. 8, pp. 1735–1780, 1997.

[13] A. Vaswani *et al.*, "Attention is all you need," in *Advances in Neural Information Processing Systems*, vol. 30, 2017, pp. 5998–6008.

[14] J. Dressel and H. Farid, "The accuracy, fairness, and limits of predicting recidivism," *Science Advances*, vol. 4, no. 1, p. eaao5580, 2018.

[15] S. Verma and J. Rubin, "Fairness definitions explained," in *Proc. Int. Workshop on Software Fairness*, 2018, pp. 1–7.

[16] J. Kleinberg, S. Mullainathan, and M. Raghavan, "Inherent trade-offs in the fair determination of risk scores," *arXiv preprint arXiv:1609.05807*, 2016.

[17] F. Kamiran and T. Calders, "Data preprocessing techniques for discrimination prevention," *Knowledge and Information Systems*, vol. 33, no. 1, pp. 1–33, 2012.

[18] M. B. Zafar, I. Valera, M. Gomez Rodriguez, and K. P. Gummadi, "Fairness constraints: Mechanisms for fair classification," in *Proc. AISTATS*, 2017, pp. 962–970.

[19] M. Hardt, E. Price, and N. Srebro, "Equality of opportunity in supervised learning," in *Advances in Neural Information Processing Systems*, vol. 29, 2016, pp. 3315–3323.

[20] R. Caruana, "Multitask learning," *Machine Learning*, vol. 28, no. 1, pp. 41–75, 1997.

[21] Y. Zhang and Q. Yang, "A survey on multi-task learning," *IEEE Trans. Knowledge and Data Engineering*, vol. 34, no. 12, pp. 5586–5609, 2021.

[22] T. Chen and C. Guestrin, "XGBoost: A scalable tree boosting system," in *Proc. 22nd ACM SIGKDD Int. Conf. on Knowledge Discovery and Data Mining (KDD '16)*, 2016, pp. 785–794.

[23] National Crime Records Bureau, "Crime in India 2017-2022," Ministry of Home Affairs, Government of India, New Delhi. [Online]. Available: https://ncrb.gov.in/. [Accessed: Jan. 2024].

[24] C. Catlett, N. Collier, and P. S. Yu, "Transportation safety and crime prediction using machine learning," *Big Data Cogn. Comput.*, vol. 4, no. 3, p. 22, 2020.

[25] Y. Wang, H. Zhang, and J. Chen, "Fairness-aware machine learning: A survey on bias detection and mitigation," *Neural Comput. Appl.*, vol. 35, no. 12, pp. 8901–8925, 2023.

[26] M. Mohammadi, A. N. Toosi, and M. P. P. N. Silva, "Algorithmic fairness in criminal justice: A systematic literature review," *J. Big Data*, vol. 10, no. 1, p. 89, 2023.

[27] X. Zhang, L. Wang, and K. He, "Deep learning for spatio-temporal crime prediction: Recent advances and challenges," *Artif. Intell. Rev.*, vol. 57, no. 2, p. 34, 2024.

[28] A. Kumar, S. Sharma, and R. Singh, "Machine learning approaches for crime analysis and prediction: A comprehensive review," *Int. J. Inf. Technol.*, vol. 14, no. 6, pp. 1205–1218, 2022.

[29] R. Gupta and P. Kumar, "Bias and fairness in AI-driven law enforcement systems: Challenges and opportunities," *J. Reliab. Intell. Environ.*, vol. 7, no. 3, pp. 189–203, 2021.

[30] S. Patel, M. Thakur, and A. Verma, "Multi-task deep learning for urban crime forecasting with demographic considerations," *Multimedia Tools Appl.*, vol. 82, no. 15, pp. 22847–22869, 2023.

[31] J. Lee, K. Park, and S. Kim, "Intersectional fairness in predictive policing: Addressing compounded vulnerabilities," *Ethics Inf. Technol.*, vol. 26, no. 1, p. 12, 2024.

[32] M. Rodriguez, A. Martinez, and F. Garcia, "Temporal attention mechanisms for crime pattern analysis," *Expert Syst. Appl.*, vol. 208, p. 118089, 2022.

[33] D. Thompson, R. Williams, and L. Brown, "Causal inference for fairness in criminal risk assessment," *J. Artif. Intell. Res.*, vol. 76, pp. 445–489, 2023.